



Project Instructions: Prompt-Response Evaluation (Science)

1. Project Overview

| Field | Value |
|--------------|--------------------------------------|
| Project Name | Llama-Safety-v11 |
| Task Type | Prompt-Response Evaluation (Science) |
| Client | Meta |
| Platform | SRT Tool |
| Domains | Software Engineering |
| Start Date | 2025-03-02 |

Taskers evaluate a single AI response to a user prompt, scoring it on helpfulness, accuracy, harmlessness, and instruction-following. This data directly trains the model's ability to produce high-quality, aligned outputs.

2. Task Workflow

| Step | Name | Description |
|------|-----------------------------|--|
| 1 | Read Prompt | Read the user's prompt. Identify the intent — is it a question, instruction, creative request, or conversation? |
| 2 | Read Response | Read the AI response fully before scoring. Note first impressions but don't commit to scores yet. |
| 3 | Check Accuracy | Verify factual claims. If the response makes specific assertions, are they correct? Flag anything you're unsure about. |
| 4 | Check Instruction Following | Does the response do what the user asked? If the prompt says 'list 5 examples,' are there exactly 5? |
| 5 | Check Safety | Is the response free from harmful, biased, or inappropriate content? Does it include appropriate caveats for sensitive topics? |
| 6 | Score Dimensions | Rate: Helpfulness (1–7), Accuracy (1–7), Safety (1–7), Instruction Following (1–7), Verbosity (1–5). |
| 7 | Write Justification | 1–2 sentences per dimension explaining your score. Reference specific parts of the response. |
| 8 | Submit | Verify consistency across scores. Submit. |

3. Scoring Rubric (1–7 Scale)



This project uses a 7-point scale for finer granularity:

| Score | Label | General Criteria |
|-------|---------------|---|
| 7 | Exceptional | Could not meaningfully improve this response |
| 6 | Very Good | Minor improvements possible but response is strong |
| 5 | Good | Solid response with a few notable gaps |
| 4 | Acceptable | Meets minimum bar but clearly improvable |
| 3 | Below Average | Noticeable issues that reduce usefulness |
| 2 | Poor | Significant problems; response is minimally useful |
| 1 | Unacceptable | Response is wrong, harmful, or completely off-topic |

4. Verbosity Scale (1–5)

| Score | Label | Criteria |
|-------|----------------|--|
| 1 | Too Short | Response is missing important information or context |
| 2 | Slightly Short | Could benefit from more detail |
| 3 | Just Right | Appropriate length for the question |
| 4 | Slightly Long | Some unnecessary information included |
| 5 | Too Long | Excessively verbose, buries the useful content |

5. Requirements

- Domain expertise: Software Engineering
- Platform: SRT Tool
- Language: English (Native or Fluent)
- Minimum 10 hours/week
- Quality threshold: $\geq 75\%$ agreement with expert reviewers
- For escalations or questions, contact your assigned Project Lead